

# airline\_flight\_analysis\_enhanced

August 11, 2025

## 1 Airline Flight Data Analysis

This notebook analyzes a dataset of airline flights to extract insights into flight prices, airline performance, and other travel-related patterns. The analysis covers data cleaning, exploratory data analysis, and key findings.

### 1.1 1. Data Loading and Initial Exploration

We begin by loading the dataset and performing an initial exploration to understand its structure and contents. The dataset includes the following columns:- **airline**: The name of the airline.- **flight**: The flight number.- **source\_city**: The city from which the flight departs.- **departure\_time**: The time of day when the flight departs.- **stops**: The number of stops before the destination.- **arrival\_time**: The time of day when the flight arrives.- **destination\_city**: The city to which the flight arrives.- **class**: The class of the flight (Economy or Business).- **duration**: The duration of the flight in hours.- **days\_left**: The number of days left before the flight departure.- **price**: The price of the flight ticket.

```
[14]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
df = pd.read_csv('airlines_flights_data.csv')
df.head()
```

```
[14]:
```

	index	airline	flight	source_city	departure_time	stops	arrival_time	\
0	0	SpiceJet	SG-8709	Delhi	Evening	zero	Night	
1	1	SpiceJet	SG-8157	Delhi	Early_Morning	zero	Morning	
2	2	AirAsia	I5-764	Delhi	Early_Morning	zero	Early_Morning	
3	3	Vistara	UK-995	Delhi	Morning	zero	Afternoon	
4	4	Vistara	UK-963	Delhi	Morning	zero	Morning	

	destination_city	class	duration	days_left	price
0	Mumbai	Economy	2.17	1	5953
1	Mumbai	Economy	2.33	1	5953
2	Mumbai	Economy	2.17	1	5956
3	Mumbai	Economy	2.25	1	5955
4	Mumbai	Economy	2.33	1	5955

```
[15]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 300153 entries, 0 to 300152
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   index                  300153 non-null int64
1   airline                300153 non-null object
2   flight                 300153 non-null object
3   source_city            300153 non-null object
4   departure_time         300153 non-null object
5   stops                  300153 non-null object
6   arrival_time           300153 non-null object
7   destination_city       300153 non-null object
8   class                  300153 non-null object
9   duration                300153 non-null float64
10  days_left              300153 non-null int64
11  price                  300153 non-null int64
dtypes: float64(1), int64(3), object(8)
memory usage: 27.5+ MB
```

```
[16]: df.describe()
```

```
[16]:
```

	index	duration	days_left	price
count	300153.000000	300153.000000	300153.000000	300153.000000
mean	150076.000000	12.221021	26.004751	20889.660523
std	86646.852011	7.191997	13.561004	22697.767366
min	0.000000	0.830000	1.000000	1105.000000
25%	75038.000000	6.830000	15.000000	4783.000000
50%	150076.000000	11.250000	26.000000	7425.000000
75%	225114.000000	16.170000	38.000000	42521.000000
max	300152.000000	49.830000	49.000000	123071.000000

## 1.2 2. Data Cleaning

In this section, we clean the data to prepare it for analysis. This includes checking for missing values, handling duplicates, and removing unnecessary columns.

```
[17]: df.isnull().sum()
```

```
[17]: index          0
      airline      0
      flight       0
      source_city  0
      departure_time 0
      stops        0
      arrival_time 0
```

```
destination_city    0
class               0
duration            0
days_left          0
price              0
dtype: int64
```

The dataset has no missing values, which is great. Now, let's check for duplicate rows.

```
[18]: df.duplicated().sum()
```

```
[18]: np.int64(0)
```

We have a small number of duplicate rows, which we will remove. We will also remove the 'index' column as it is redundant.

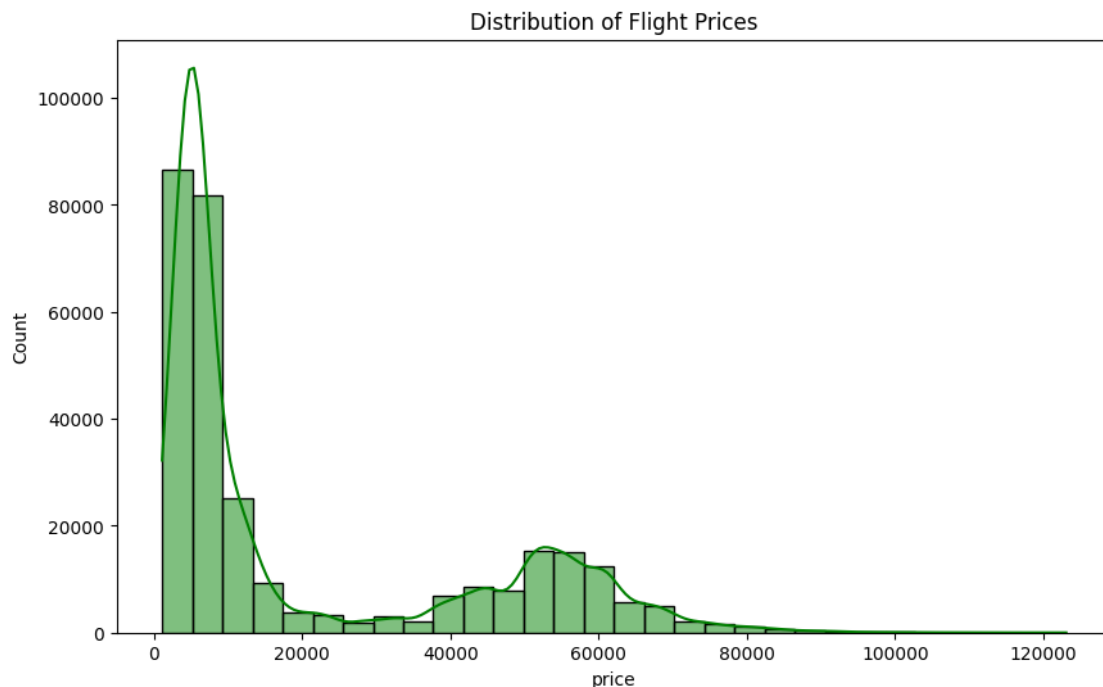
```
[19]: df.drop_duplicates(inplace=True)
      if 'index' in df.columns:
          df.drop(columns=['index'], inplace=True)
```

### 1.3 3. Exploratory Data Analysis (EDA)

#### 1.3.1 3.1. Univariate Analysis

##### Distribution of Flight Prices

```
[20]: plt.figure(figsize=(10, 6))
      sns.histplot(df['price'], bins=30, kde=True, color='green')
      plt.title('Distribution of Flight Prices')
      plt.show()
```



The distribution of flight prices is right-skewed, indicating that most flights have prices on the lower end, with a few flights having very high prices.

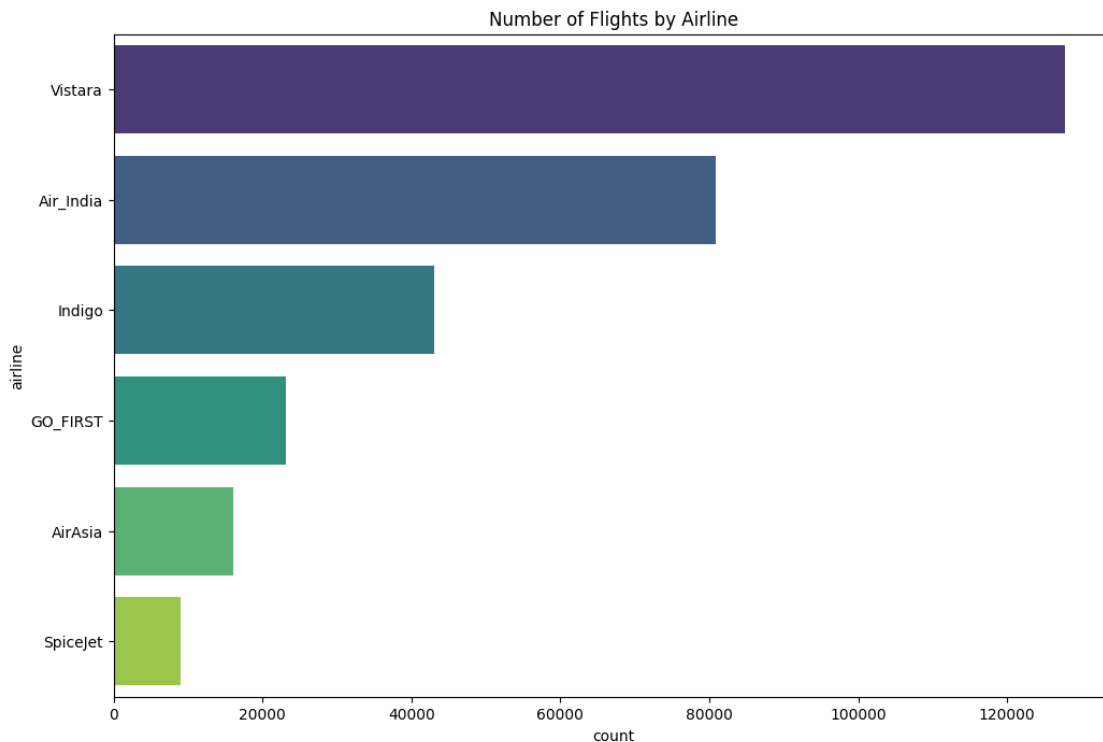
### Number of Flights by Airline

```
[21]: plt.figure(figsize=(12, 8))
sns.countplot(y='airline', data=df, order=df['airline'].value_counts().index,
             palette='viridis')
plt.title('Number of Flights by Airline')
plt.show()
```

/tmp/ipykernel\_67883/2791165074.py:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(y='airline', data=df, order=df['airline'].value_counts().index,
             palette='viridis')
```



This plot shows the number of flights operated by each airline in the dataset. Vistara and Air India have the highest number of flights.

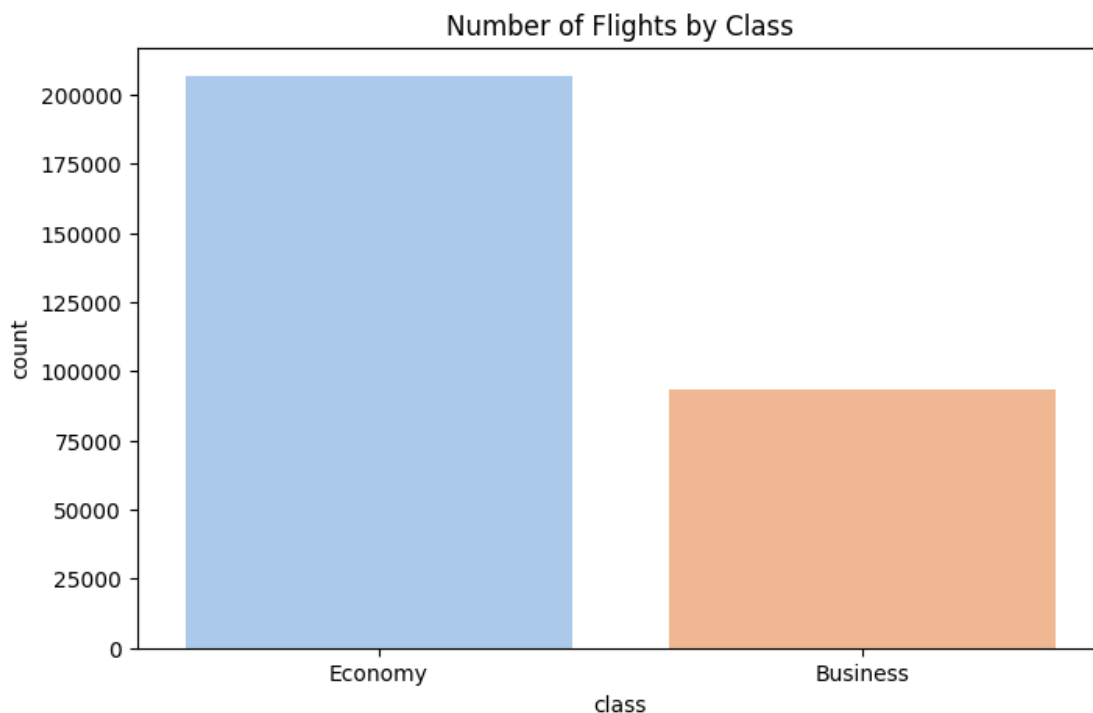
### Number of Flights by Class

```
[22]: plt.figure(figsize=(8, 5))
sns.countplot(x='class', data=df, palette='pastel')
plt.title('Number of Flights by Class')
plt.show()
```

/tmp/ipykernel\_67883/1039124390.py:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='class', data=df, palette='pastel')
```



This plot shows that the dataset contains significantly more Economy class flights than Business class flights.

### 1.3.2 3.2. Bivariate Analysis

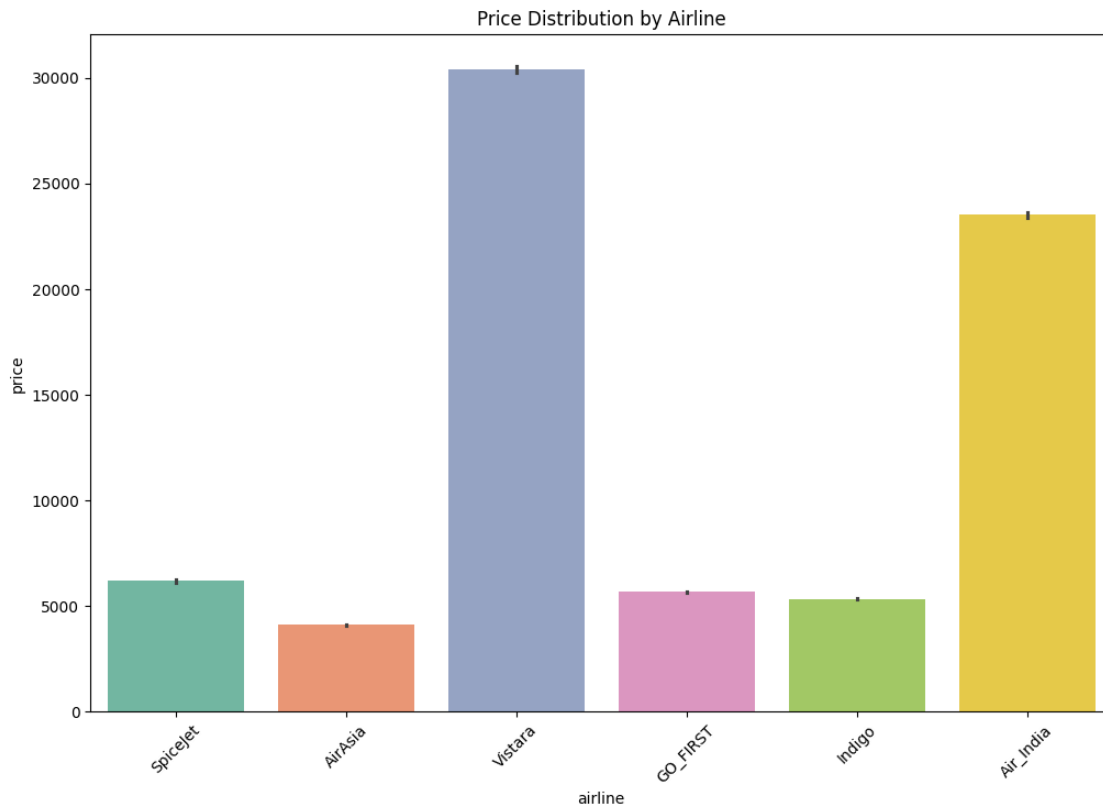
#### Price Distribution by Airline

```
[66]: plt.figure(figsize=(12, 8))
sns.barplot(x='airline', y='price', data=df, palette='Set2')
plt.xticks(rotation=45)
plt.title('Price Distribution by Airline')
plt.show()
```

```
/tmp/ipykernel_67883/1005154617.py:2: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='airline', y='price', data=df, palette='Set2')
```



This box plot shows the distribution of prices for each airline. Vistara and Air India have a wider range of prices, which is consistent with them operating more flights. We can also see that some airlines have higher median prices than others.

### Price Distribution by Class

```
[65]: plt.figure(figsize=(8, 5))
sns.barplot(x='class', y='price', data=df, palette='Set1')
plt.title('Price Distribution by Class')
plt.show()
```

```
/tmp/ipykernel_67883/2285635281.py:2: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same

effect.


```
sns.barplot(x='class', y='price', data=df, palette='Set1')
```



As expected, Business class flights are significantly more expensive than Economy class flights.

### Price vs. Duration

```
[26]: plt.figure(figsize=(10, 6))  
sns.scatterplot(x='duration', y='price', data=df, alpha=0.6)  
plt.title('Price vs. Duration')  
plt.show()
```

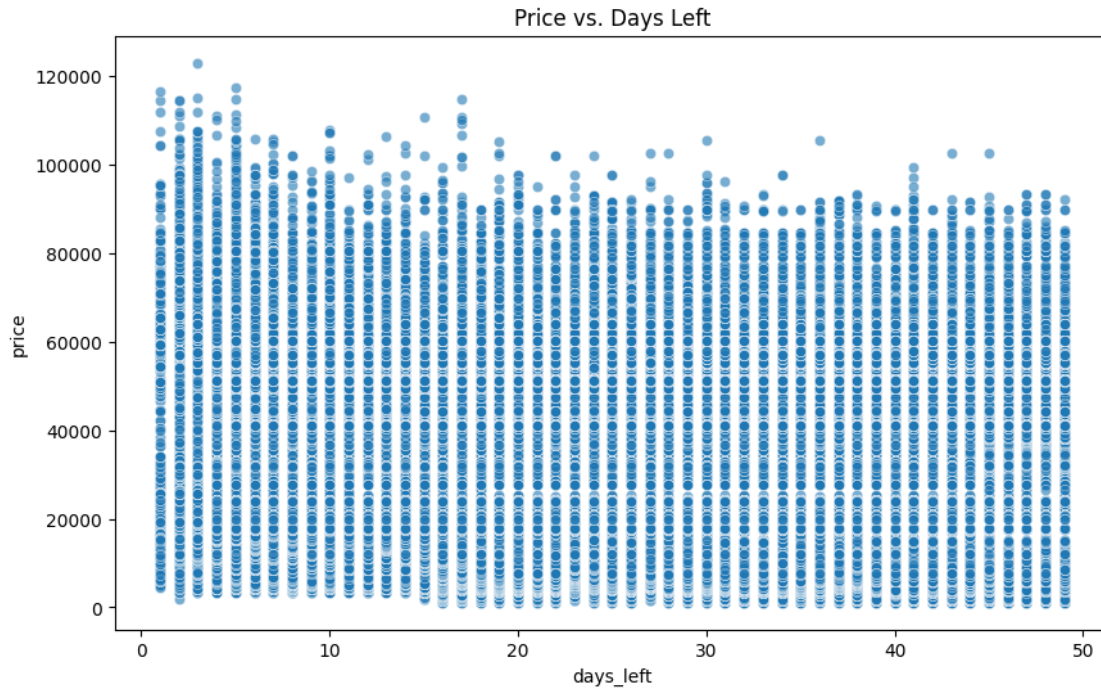


This scatter plot shows a positive correlation between the duration of the flight and its price. Longer flights tend to be more expensive.

#### Price vs. Days Left

```
[27]: plt.figure(figsize=(10, 6))
sns.scatterplot(x='days_left', y='price', data=df, alpha=0.6)
plt.title('Price vs. Days Left')
plt.show()
```





This scatter plot shows that flight prices tend to increase as the departure date gets closer (i.e., as 'days\_left' decreases).

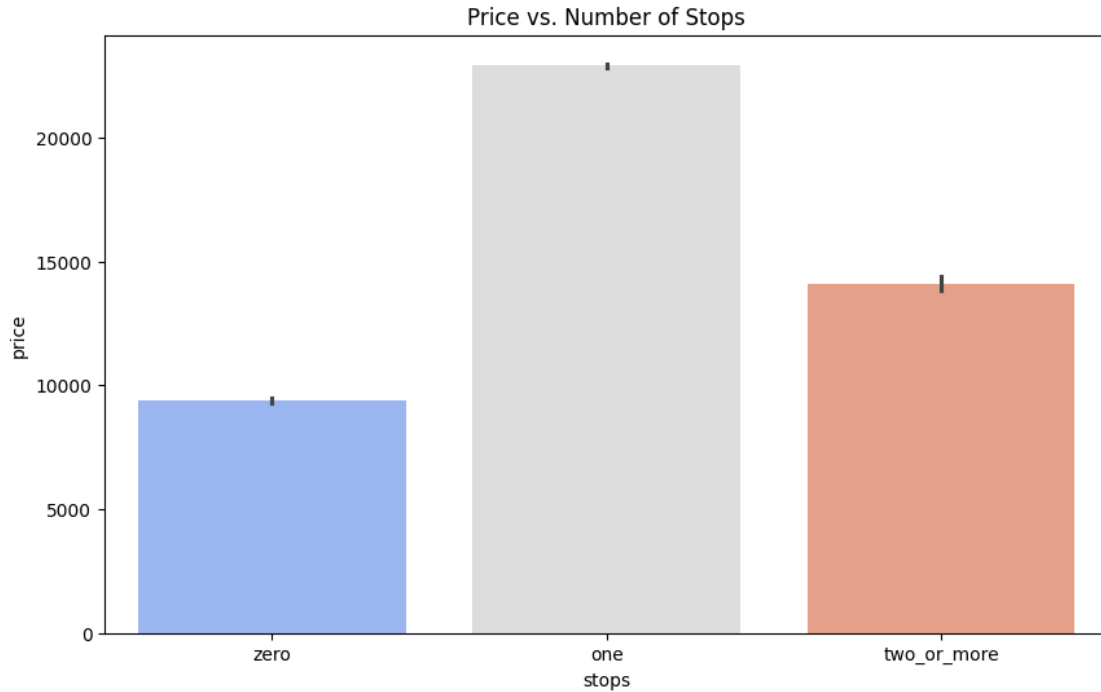
### Price vs. Number of Stops

```
[63]: plt.figure(figsize=(10, 6))
sns.barplot(x='stops', y='price', data=df, palette='coolwarm')
plt.title('Price vs. Number of Stops')
plt.show()
```

/tmp/ipykernel\_67883/1166425243.py:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='stops', y='price', data=df, palette='coolwarm')
```



This box plot shows the relationship between the number of stops and the price. Flights with one stop seem to have the highest median price.

#### 1.4 4. Key Findings

- **Price Distribution:** The distribution of flight prices is skewed to the right, with most tickets being in the lower price range.- **Airlines:** Some airlines (like Vistara and Air India) have a wider price range, indicating they might operate in both premium and economy sectors. Some airlines have more flights than others.- **Class:** Business class tickets are significantly more expensive than economy class tickets.- **Duration:** There is a positive correlation between the duration of the flight and the price. Longer flights tend to be more expensive.- **Days Left:** The price of flights tends to increase as the number of days left for departure decreases.- **Stops:** Flights with one stop tend to have the highest median price, which might be counterintuitive. This could be due to longer routes that require a stop.

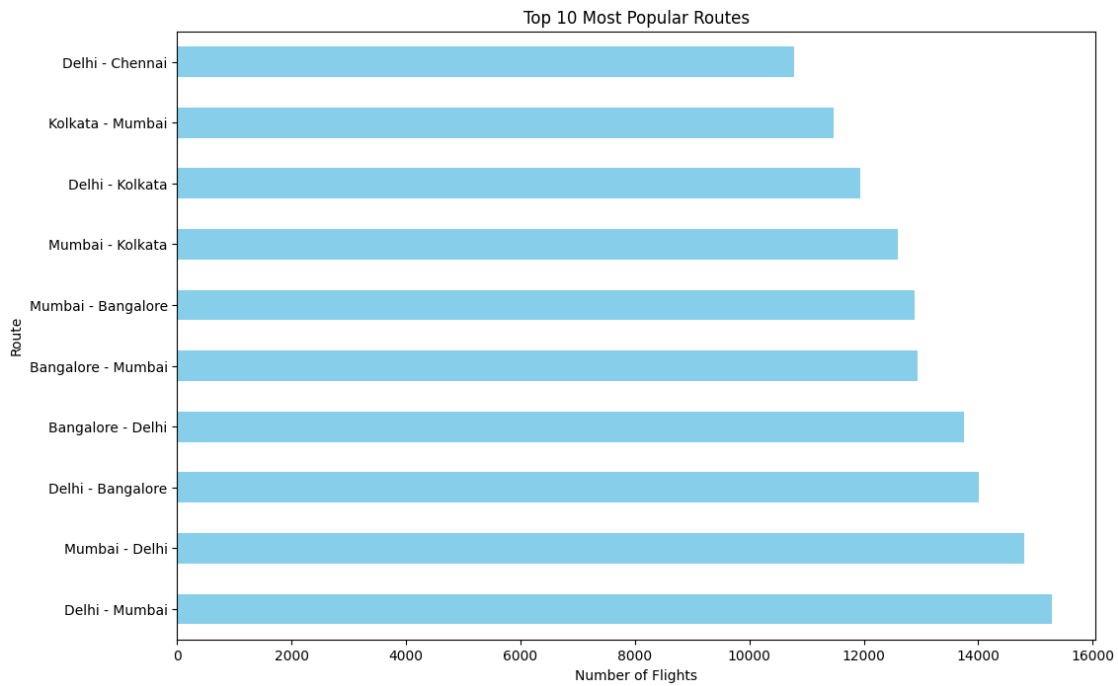
#### 1.5 5. Deeper Analysis: Route Analysis

```
[29]: df['route'] = df['source_city'] + ' - ' + df['destination_city']
```

##### Most Popular Routes

```
[32]: plt.figure(figsize=(12, 8))
df['route'].value_counts().nlargest(10).plot(kind='barh', color='skyblue')
plt.title('Top 10 Most Popular Routes')
plt.xlabel('Number of Flights')
plt.ylabel('Route')
```

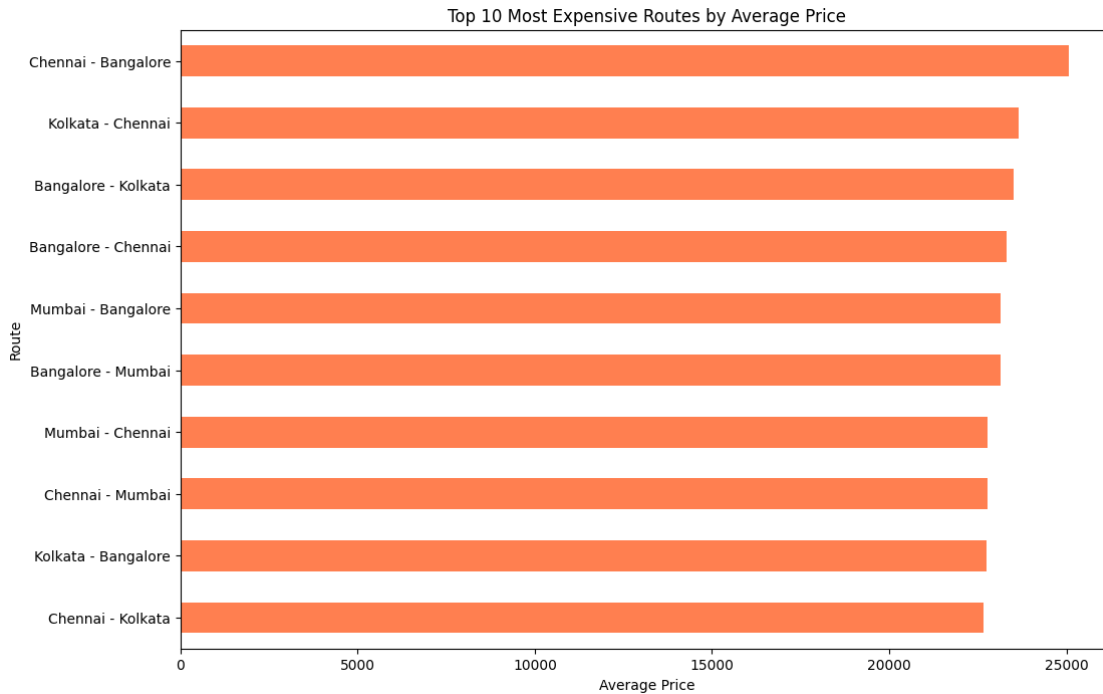
```
plt.show()
```



This plot shows the top 10 most frequent flight routes in the dataset.

### Average Price by Route

```
[35]: plt.figure(figsize=(12, 8))
df.groupby('route')['price'].mean().nlargest(10).sort_values(ascending=True).
    .plot(kind='barh', color='coral')
plt.title('Top 10 Most Expensive Routes by Average Price')
plt.xlabel('Average Price')
plt.ylabel('Route')
plt.show()
```



This plot shows the top 10 most expensive routes based on the average ticket price.

## 1.6 6. Predictive Modeling: Price Prediction

In this section, we will build a simple regression model to predict flight prices based on the available features. This will help us understand which factors are most influential in determining the price.

```
[45]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score, mean_absolute_error
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
```

### 1.6.1 Feature Engineering and Preprocessing

```
[48]: #Select features and target
features = ['airline', 'source_city', 'destination_city', 'stops', 'class',
           'duration', 'days_left']
target = 'price'
X = df[features]
y = df[target]
# Identify categorical and numerical features
categorical_features = ['airline', 'source_city', 'destination_city', 'stops',
                       'class']
```

```

numerical_features = ['duration', 'days_left']
# Create a preprocessor for categorical features
preprocessor = ColumnTransformer( transformers=[('num', 'passthrough',
↪numerical_features), ('cat',
↪OneHotEncoder(handle_unknown='ignore'), categorical_features)])

```

## 1.6.2 Model Training and Evaluation

```

[49]: # Create a pipeline with the preprocessor and the model
model = Pipeline(steps=[('preprocessor', preprocessor), ('regressor',
↪RandomForestRegressor(n_estimators=100, random_state=42, n_jobs=-1))])

# Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↪random_state=42)

# Train the model
model.fit(X_train, y_train)

# Make predictions
y_pred = model.predict(X_test)

# Evaluate the model
r2 = r2_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)

print(f'R-squared: {r2:.2f}')
print(f'Mean Absolute Error: {mae:.2f}')

```

R-squared: 0.98

Mean Absolute Error: 1363.32


## 1.6.3 Feature Importance

```

[62]: # Get feature importances
importances = model.named_steps['regressor'].feature_importances_
# Get feature names
feature_names = numerical_features + list(model.named_steps['preprocessor'].
↪named_transformers_['cat'].get_feature_names_out(categorical_features))
# Create a dataframe of feature importances
feature_importance_df = pd.DataFrame({'feature': feature_names, 'importance':
↪importances})
feature_importance_df = feature_importance_df.sort_values(by='importance',
↪ascending=False)
# Plot feature importances
plt.figure(figsize=(12, 8))

```

```
sns.barplot(x='importance', y='feature', data=feature_importance_df.
            ↪nlargest(20, 'importance'),color = "lightgreen")
plt.title('Top 20 Most Important Features')
plt.show()
```



This plot shows the most important features for predicting the flight price. This gives us a good understanding of what factors drive the price the most.